

Airbnb Dynamic Pricing Recommendation Engine — Full Package

Generated: Detailed deliverables including full code, API server, Tableau fields, and PDF generator.

1. Overview

This document contains the complete code, deployment instructions, Tableau calculated fields, and PDF report generator for the Airbnb Dynamic Pricing Recommendation Engine. Use the code blocks as copy-paste-ready files. Follow 'How to run' sections for local execution.

2. pricing_engine.py (Full script)

```
# pricing_engine.py
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.metrics import mean_squared_error, mean_absolute_error
from xgboost import XGBRegressor
import joblib
import os
from datetime import datetime

RANDOM_STATE = 42

def load_data(path):
    df = pd.read_csv(path, parse_dates=['date'], dayfirst=True, low_memory=False)
    return df

def clean_basic(df):
    def to_num(x):
        if pd.isna(x): return np.nan
        s = str(x).replace('$','').replace(',','').strip()
        return float(s) if s!="" else np.nan
    for c in ['price', 'cleaning_fee', 'security_deposit', 'competitor_price_avg']:
        if c in df.columns:
            df[c] = df[c].astype(str).replace('nan','', regex=False).apply(lambda v: np.nan if v=='' else to_num(v))
    if 'amenities' in df.columns:
        df['amenities_count'] = df['amenities'].fillna('').apply(lambda s: s.count(',')+1 if s.strip() else 0)
    if 'host_since' in df.columns:
        df['host_since'] = pd.to_datetime(df['host_since'], errors='coerce')
        df['host_tenure_days'] = (pd.Timestamp.now() - df['host_since']).dt.days.fillna(0)
    if 'availability_365' in df.columns:
        df['occupancy_estimate'] = 1.0 - (df['availability_365'].fillna(365) / 365.0)
    if 'date' in df.columns:
        df['month'] = df['date'].dt.month
        df['dayofweek'] = df['date'].dt.dayofweek
        df['is_weekend'] = df['dayofweek'].isin([5,6]).astype(int)
    numeric_cols = df.select_dtypes(include=['float64', 'int64']).columns.tolist()
    for c in numeric_cols:
        df[c] = pd.to_numeric(df[c], errors='coerce')
    return df

def build_pipeline(categorical_cols, numeric_cols):
    cat_pipe = Pipeline([
        ('impute', SimpleImputer(strategy='constant', fill_value='unknown')),
        ('ohe', OneHotEncoder(handle_unknown='ignore', sparse=False))
    ])
    num_pipe = Pipeline([
        ('impute', SimpleImputer(strategy='median')),
        ('scale', StandardScaler())
    ])
    pre = ColumnTransformer([
        ('cat', cat_pipe, categorical_cols),
        ('num', num_pipe, numeric_cols),
    ], remainder='drop')
    # Store attributes for later use
    pre.categorical_features = categorical_cols
    pre.numeric_features = numeric_cols
    return pre

def train_model(df, model_path='model.joblib', preprocess_path='preprocessor.joblib'):
```

```

df = df.copy()
df = df[df['price'].notna()]
df['log_price'] = np.log1p(df['price'])
categorical_cols = ['city', 'neighbourhood', 'property_type', 'room_type']
categorical_cols = [c for c in categorical_cols if c in df.columns]
numeric_cols = ['accommodates', 'bedrooms', 'beds', 'bathrooms',
                 'amenities_count', 'review_scores_rating', 'number_of_reviews',
                 'reviews_per_month', 'host_tenure_days', 'occupancy_estimate',
                 'month', 'dayofweek', 'is_weekend', 'competitor_price_avg']
numeric_cols = [c for c in numeric_cols if c in df.columns]
X = df[categorical_cols + numeric_cols].copy()
y = df['log_price'].values
pre = build_pipeline(categorical_cols, numeric_cols)
if 'date' in df.columns:
    df_sorted = df.sort_values('date')
    split_index = int(len(df_sorted) * 0.8)
    train_idx = df_sorted.index[:split_index]
    test_idx = df_sorted.index[split_index:]
    X_train, X_test = X.loc[train_idx], X.loc[test_idx]
    y_train, y_test = y[train_idx], y[test_idx]
else:
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=RANDOM_STATE)
X_train_t = pre.fit_transform(X_train)
X_test_t = pre.transform(X_test)
xgb = XGBRegressor(objective='reg:squarederror', random_state=RANDOM_STATE, n_jobs=4)
param_grid = {'n_estimators':[100,200], 'max_depth':[4,6], 'learning_rate':[0.05,0.1]}
grid = GridSearchCV(xgb, param_grid, cv=3, scoring='neg_root_mean_squared_error', verbose=1, n_jobs=-1)
grid.fit(X_train_t, y_train)
best = grid.best_estimator_
preds_log = best.predict(X_test_t)
preds_price = np.expm1(preds_log)
y_test_price = np.expm1(y_test)
rmse = mean_squared_error(y_test_price, preds_price, squared=False)
mae = mean_absolute_error(y_test_price, preds_price)
print(f"Model trained. Test RMSE (raw price): {rmse:.2f}, MAE: {mae:.2f}")
joblib.dump(best, model_path)
joblib.dump(pre, preprocess_path)
return {
    'model': best,
    'preprocessor': pre,
    'categorical_cols': categorical_cols,
    'numeric_cols': numeric_cols,
    'rmse': rmse,
    'mae': mae
}

def predict_price_for_df(df_in, artifacts):
    pre = artifacts['preprocessor']
    model = artifacts['model']
    X = df_in[artifacts['categorical_cols']] + artifacts['numeric_cols'].copy()
    X_t = pre.transform(X)
    preds_log = model.predict(X_t)
    preds_price = np.expm1(preds_log)
    return preds_price

def suggest_price(row, predicted_price, slider_percent=0.0, min_price_col=None, max_multiplier=1.6):
    base = predicted_price * (1.0 + slider_percent)
    floors = []
    if min_price_col and min_price_col in row and pd.notna(row[min_price_col]):
        floors.append(float(row[min_price_col]))
    if 'price' in row and pd.notna(row['price']):
        floors.append(float(row['price'])*0.5)
    floors.append(predicted_price*0.6)
    min_price = max(floors) if floors else predicted_price*0.6
    max_price = predicted_price * max_multiplier
    if 'review_scores_rating' in row and pd.notna(row['review_scores_rating']):
        if row['review_scores_rating'] < 75 and base > predicted_price * 1.1:
            base = predicted_price * 1.1
    suggested = float(np.clip(base, min_price, max_price))
    return round(suggested, 2)

def run_full_training(csv_path, artifacts_dir='artifacts'):
    os.makedirs(artifacts_dir, exist_ok=True)
    df = load_data(csv_path)
    df = clean_basic(df)
    artifacts = train_model(df, model_path=os.path.join(artifacts_dir, 'model.joblib'),
                           preprocess_path=os.path.join(artifacts_dir, 'preproc.joblib'))
    predicted_prices = predict_price_for_df(df, artifacts)
    df['predicted_price'] = np.round(predicted_prices, 2)
    df['suggested_price'] = df.apply(lambda r: suggest_price(r, r['predicted_price'], slider_percent=0.0),
                                     axis=1)
    out_csv = os.path.join(artifacts_dir, 'price_suggestions.csv')

```

```
df.to_csv(out_csv, index=False)
print(f"Saved suggestions to {out_csv}")
return artifacts, df

if __name__ == '__main__':
    import argparse
    parser = argparse.ArgumentParser(description="Train Airbnb Pricing Engine")
    parser.add_argument('--data', required=True, help='path to historical data CSV')
    parser.add_argument('--out', default='artifacts', help='output artifacts folder')
    args = parser.parse_args()
    artifacts, df = run_full_training(args.data, artifacts_dir=args.out)
```

3. FastAPI server (app.py)

```
# app.py
from fastapi import FastAPI, HTTPException
from pydantic import BaseModel, Field
from typing import List, Optional
import joblib
import numpy as np
import pandas as pd
import uvicorn
import os

ARTIFACTS_DIR = os.environ.get('ARTIFACTS_DIR', 'artifacts')
MODEL_PATH = os.path.join(ARTIFACTS_DIR, 'model.joblib')
PREPROC_PATH = os.path.join(ARTIFACTS_DIR, 'preproc.joblib')

model = joblib.load(MODEL_PATH)
preproc = joblib.load(PREPROC_PATH)

app = FastAPI(title="Airbnb Pricing Engine API")

class PredictRequest(BaseModel):
    listing_id: Optional[str]
    city: Optional[str] = None
    neighbourhood: Optional[str] = None
    property_type: Optional[str] = None
    room_type: Optional[str] = None
    accommodates: Optional[float] = None
    bedrooms: Optional[float] = None
    beds: Optional[float] = None
    bathrooms: Optional[float] = None
    amenities_count: Optional[float] = None
    review_scores_rating: Optional[float] = None
    number_of_reviews: Optional[float] = None
    reviews_per_month: Optional[float] = None
    host_tenure_days: Optional[float] = None
    occupancy_estimate: Optional[float] = None
    month: Optional[int] = None
    dayofweek: Optional[int] = None
    is_weekend: Optional[int] = None
    competitor_price_avg: Optional[float] = None
    price: Optional[float] = None
    min_nightly_price: Optional[float] = None
    slider_percent: Optional[float] = Field(0.0, description="e.g. 0.1 for +10%, -0.2 for -20%")

class BatchRequest(BaseModel):
    rows: List[PredictRequest]

def _to_dataframe(req: PredictRequest):
    d = req.dict()
    return pd.DataFrame([d])

def _predict_from_df(df: pd.DataFrame):
    try:
        X_t = preproc.transform(df)
    except Exception as e:
        raise HTTPException(status_code=400, detail=f"Preprocessor transform failed: {e}")
    preds_log = model.predict(X_t)
    preds = np.expm1(preds_log)
    return preds

def suggest_price_from_row(row: pd.Series, predicted_price: float, slider_percent: float=0.0, max_multiplier: float=1.1):
    base = predicted_price * (1.0 + slider_percent)
    floors = []
    if 'min_nightly_price' in row and pd.notna(row.get('min_nightly_price')):
        floors.append(float(row['min_nightly_price']))
    if 'price' in row and pd.notna(row.get('price')):
        floors.append(float(row['price']) * 0.5)
    floors.append(predicted_price * 0.6)
    min_price = max(floors) if floors else predicted_price * 0.6
    max_price = predicted_price * max_multiplier
    if 'review_scores_rating' in row and pd.notna(row.get('review_scores_rating')):
        if row['review_scores_rating'] < 75 and base > predicted_price * 1.1:
            base = predicted_price * 1.1
    suggested = float(np.clip(base, min_price, max_price))
    return round(suggested, 2)

@app.post("/predict")
def predict(req: PredictRequest):
    df = _to_dataframe(req)
```

```
preds = _predict_from_df(df)
predicted_price = float(preds[0])
suggested = suggest_price_from_row(df.iloc[0], predicted_price, slider_percent=req.slider_percent or None)
return {
    "listing_id": req.listing_id,
    "predicted_price": round(predicted_price, 2),
    "suggested_price": suggested,
    "slider_percent": req.slider_percent
}

@app.post("/batch_predict")
def batch_predict(req: BatchRequest):
    rows = [r.dict() for r in req.rows]
    df = pd.DataFrame(rows)
    preds = _predict_from_df(df)
    results = []
    for i, r in df.iterrows():
        p = float(preds[i])
        slider = float(r.get('slider_percent', 0.0) or 0.0)
        suggested = suggest_price_from_row(r, p, slider_percent=slider)
        results.append({
            "index": int(i),
            "predicted_price": round(p, 2),
            "suggested_price": suggested,
            "slider_percent": slider
        })
    return {"results": results}

if __name__ == "__main__":
    uvicorn.run("app:app", host="0.0.0.0", port=8000, reload=True)
```

4. Tableau: Parameters & Calculated Fields

Parameter: price_adjustment - Data type: Float - Current value: 0.0 - Range: Min -0.20, Max 0.30, Step 0.01
Calculated Field: Suggested Price (raw) [Predicted_Price] * (1 + [price_adjustment])
Calculated Field: Suggested Price (clipped) FLOAT(IF [Review_Scores_Rating] < 75 AND [Suggested Price (raw)] > [Predicted_Price] * 1.1 THEN MIN([Predicted_Price] * 1.1, [Predicted_Price] * 1.6) ELSE MIN([Predicted_Price] * 1.6, MAX(IFNULL([min_nightly_price], 0), [Predicted_Price] * 0.6, IFNULL([price],0) * 0.5, [Suggested Price (raw)])) END)
Calculated Field: Opportunity % IFNULL([price],0) = 0 THEN 0 ELSE ([Suggested Price (clipped)] - [price]) / [price] END

5. PDF Report Generator (generate_report.py)

```
# generate_report.py
import pandas as pd
import matplotlib.pyplot as plt
from reportlab.lib.pagesizes import A4
from reportlab.pdfgen import canvas
from reportlab.lib.utils import ImageReader
import os

def plot_top_cities(df, outpath='top_cities.png'):
    agg = df.groupby('city').agg({'price':'mean','predicted_price':'mean','suggested_price':'mean'}).sort_values(by='suggested_price', ascending=False)
    ax = agg[['price','predicted_price','suggested_price']].plot(kind='bar', figsize=(10,4))
    ax.set_title("Top 10 Cities - Avg Prices")
    plt.tight_layout()
    plt.savefig(outpath)
    plt.close()
    return outpath

def plot_residuals(df, outpath='residuals.png'):
    df = df[df['price'].notna()]
    df['residual'] = df['predicted_price'] - df['price']
    ax = df['residual'].hist(bins=50, figsize=(8,4))
    ax.set_title('Residuals (Predicted - Actual)')
    plt.tight_layout()
    plt.savefig(outpath)
    plt.close()
    return outpath

def create_pdf_report(csv_path='artifacts/price_suggestions.csv', pdf_path='price_report.pdf'):
    df = pd.read_csv(csv_path)
    city_plot = plot_top_cities(df)
    resid_plot = plot_residuals(df)

    c = canvas.Canvas(pdf_path, pagesize=A4)
    width, height = A4

    c.setFont("Helvetica-Bold", 20)
    c.drawCentreString(width/2, height-100, "Airbnb Pricing Suggestions Report")
    c.setFont("Helvetica", 12)
    c.drawString(50, height-130, f"Generated: {pd.Timestamp.now().strftime('%Y-%m-%d %H:%M')}")
    c.drawString(50, height-150, f"Total listings analyzed: {len(df)}")
    c.showPage()

    c.setFont("Helvetica-Bold", 16)
    c.drawString(50, height-60, "Top Cities: Average Prices")
    img = ImageReader(city_plot)
    c.drawImage(img, 50, height-500, width=500, height=350, preserveAspectRatio=True)
    c.showPage()

    c.setFont("Helvetica-Bold", 16)
    c.drawString(50, height-60, "Model Residuals")
    img = ImageReader(resid_plot)
    c.drawImage(img, 50, height-500, width=500, height=350, preserveAspectRatio=True)
    c.showPage()

    c.setFont("Helvetica-Bold", 16)
    c.drawString(50, height-60, "Top 10 Upsell Opportunities (Suggested > Current)")
    opportunities = df[df['price'].notna()].copy()
    opportunities['opp_pct'] = (opportunities['suggested_price'] - opportunities['price']) / opportunities['price']
    opportunities = opportunities.sort_values('opp_pct', ascending=False).head(10)
    y = height - 100
    c.setFont("Helvetica", 10)
    for _, row in opportunities.iterrows():
        line = f"{str(row.get('listing_id', ''))[:15]:15} | {str(row.get('city', ''))[:15]:15} | Current: {row['current_price']:.2f} | Suggested: {row['suggested_price']:.2f} | Opp. Pct: {row['opp_pct']:.2f}"
        c.drawString(50, y, line)
        y -= 14
        if y < 80:
            c.showPage()
            y = height - 80
    c.save()
    print(f"Saved PDF: {pdf_path}")
    return pdf_path

if __name__ == '__main__':
    create_pdf_report()
```

6. How to run everything locally

1) Install dependencies: pip install pandas numpy scikit-learn xgboost joblib matplotlib reportlab fastapi uvicorn
2) Train model and create suggestions CSV: python pricing_engine.py --data path/to/historical_airbnb.csv --out artifacts
3) Start FastAPI: Ensure artifacts/model.joblib and artifacts/preproc.joblib are present. uvicorn app:app --host 0.0.0.0 --port 8000
4) Generate PDF report: python generate_report.py